Journal of Statistical Software

MMMMMM YYYY, Volume VV, Issue II. *doi: 10.18637/jss.v000.i00*

bioLeak: Leakage-Aware Modeling and Diagnostics for Machine Learning in R

Selçuk Korkmaz
Trakya University

Abstract

Data leakage remains a recurrent source of optimistic bias in biomedical machine learning studies. Standard row-wise cross-validation and globally estimated preprocessing steps are often inappropriate for data with repeated measurements, study-level heterogeneity, batch effects, or temporal dependencies. This paper describes `bioLeak`, an R package for constructing leakage-aware resampling workflows and for auditing fitted models for common leakage mechanisms. The package provides leakage-aware split construction, train-fold-only preprocessing, cross-validated model fitting, nested hyperparameter tuning, post hoc leakage audits, and HTML reporting. Primary interfaces include `make_split_plan()`, `fit_resample()`, `tune_resample()`, `audit_leakage()`, `audit_report()`, and `delta_lsi()`. The implementation supports binary classification, multiclass classification, continuous regression, and survival analysis for modeling, with task-specific metrics and S4 containers for splits, fits, audits, and inflation summaries. Results reported in this manuscript are drawn from stored simulation and case-study outputs in the package paper directory rather than from newly executed experiments. The simulation artifacts show how apparent performance changes under controlled leakage mechanisms, and the case study illustrates how guarded and naive pipelines can yield materially different conclusions on multi-study transcriptomic data. The emphasis throughout is on software design, reproducible workflows, and interpretation of diagnostic output.

Keywords: biomedical machine learning, data leakage, cross-validation, reproducibility, resampling, auditing, R.

1. Introduction

Cross-validated performance estimates are often treated as evidence that a predictive model will generalize to future biomedical data. That interpretation depends on assumptions that are frequently violated in practice: observations must be exchangeable with respect to the

resampling design, preprocessing must be estimated without access to evaluation data, and the outcome must not be recoverable from proxy variables that are unavailable at prediction time. Violations of these assumptions are usually described as data leakage. Leakage is not a narrow implementation detail. It affects whether model accuracy, feature importance, calibration, and downstream biological interpretation are credible. This concern has been emphasized across data mining, statistical learning, and biomedical prediction, including early discussions of selection bias in microarray studies (Ambroise and McLachlan 2002), broader treatments of leakage in data mining (Kaufman, Rosset, Perlich, and Stitelman 2012), warnings about genomic prediction pitfalls (Simon, Radmacher, Dobbin, and McShane 2003; Subramanian and Simon 2013), and recent accounts linking leakage to wider reproducibility failures in machine-learning-based science (Kapoor and Narayanan 2023). Related evaluation biases also arise when model selection is not nested appropriately within resampling, a point stressed by Varma and Simon (2006) and Cawley and Talbot (2010).

The problem is especially acute in biomedical applications. A single study may contain repeated measurements from the same participant, technical replicates, laboratory batches, study-site effects, or longitudinal samples collected over time. In such settings, row-wise cross-validation can place correlated samples in both training and test folds. Similarly, global imputation, normalization, or feature filtering performed before resampling can transmit test-fold information into the training pipeline. These choices can inflate performance while leaving no obvious trace in a conventional model summary, a practical problem in health-care machine learning where evaluation, implementation, and deployment constraints are already difficult to align (Beam and Kohane 2018; Sendak, Elish, Gao, Futoma, Ratliff, Nichols, Bedoya, Balu, and O'Brien 2020).

Recent methodological work has documented the scale of the issue. Kapoor and Narayanan (2023) described leakage as a contributor to broader reproducibility failures in machine-learning-based science. Rosenblatt, Ben Simhon, Gal, and Gilron (2024) showed in neuroimaging that feature-selection and repeated-subject leakage can substantially inflate predictive performance, while van de Mortel and van Wingen (2025) reported that leakage-compatible practices were common enough to affect meta-analytic conclusions in psychiatric imaging. In clinical prediction, Davis, Greevy, Fonnesbeck, Lasko, Walsh, and Matheny (2024) discussed label leakage induced by workflow artifacts rather than by explicit coding mistakes. These studies do not imply that leakage can be removed by a single rule. Rather, they show that leakage is heterogeneous, domain specific, and tightly linked to evaluation design.

Existing modeling frameworks provide much of the infrastructure needed for modern predictive analysis, but they do not remove the need to define appropriate leakage boundaries. In R (R Core Team 2024), `caret`, `tidymodels`, and `mlr3` support flexible resampling, preprocessing, and model-fitting workflows (Kuhn 2008; Kuhn and Wickham 2020; Lang, Binder, Richter, Schratz, Pfisterer, Coors, Au, Casalicchio, Kotthoff, and Bischl 2019; Bischl, Sonabend, Kotthoff, and Lang 2024). These frameworks can be used safely, but they rely on the analyst to specify resampling units, grouping variables, and fold-wise preprocessing correctly. In biomedical applications, where rows are often nested within subjects, batches, studies, or time windows, the critical difficulty is not only fitting a model but encoding the structured dependencies that determine which observations may legitimately share information.

`bioLeak` was developed to support two related tasks. The first is prevention: constructing resampling and preprocessing workflows that reduce the risk of train-test contamination. The second is diagnosis: auditing fitted cross-validated models for evidence that apparent

performance may depend on leakage or confounding. The package therefore combines leakage-aware split plans, guarded preprocessing, cross-validated fitting, nested tuning, post-fitting audits, and reporting utilities in a single R workflow. The package is designed for practical use with tabular biomedical data and `SummarizedExperiment` objects, and it interoperates with common R modeling infrastructure such as `parsnip`, `recipes`, `workflows`, `yardstick`, and `rsample`.

This manuscript follows the software-oriented emphasis expected for the Journal of Statistical Software. Section 2 summarizes the leakage mechanisms that motivate the package. Section 3 describes the main software components. Section 4 presents a reproducible workflow based on the package API. Section 5 defines the `Delta Leakage Sensitivity Index` framework used to quantify inflation between naive and guarded pipelines. Sections 6 and 7 describe simulation and case-study results using stored paper artifacts. Section 8 discusses implementation details, Section 9 reviews limitations and future work, Section 10 summarizes the contribution, and Section 11 records current software availability.

2. Data Leakage in Biomedical Machine Learning

2.1. Preprocessing leakage

Preprocessing leakage occurs when statistics used to transform predictors are estimated using all samples rather than only the training fold. Common examples include global missing-value imputation, normalization, variance filtering, and feature selection performed before resampling. Moscovich and Rosset (2022) showed that even unsupervised preprocessing can bias cross-validation when it is estimated outside the resampling loop, a common failure mode in biomedical pipelines that treat normalization and filtering as preliminary cleaning rather than as model components. In genomics, the same concern appears as feature- selection bias: selecting genes before resampling can produce highly optimistic error estimates even when the downstream classifier is simple (Ambroise and McLachlan 2002; Simon *et al.* 2003; Subramanian and Simon 2013).

2.2. Dependent samples and subject overlap

Many biomedical datasets contain repeated measures, multi-region samples from the same patient, family-related observations, or longitudinal observations from the same person. If row-wise resampling is applied, records from the same latent unit can appear in both training and test folds. In that case the model may recognize subject identity rather than the biological signal of interest. This problem is especially acute when outcomes are stable within subject or when repeated measurements are dense. More generally, the problem belongs to the broader class of structured-resampling failures discussed by Roberts, Bahn, Ciuti, Boyce, Elith, Guillera-Arroita, Hauenstein, Lahoz-Monfort, Schröder, Thuiller, Warton, Wintle, Hartig, and Dormann (2017), Wenger and Olden (2012), and Valavi, Elith, Lahoz-Monfort, and Guillera-Arroita (2019), where naive row-wise fold assignment ignores dependencies induced by space, time, hierarchy, or shared latent units.

2.3. Batch and study confounding

High-throughput molecular assays frequently exhibit batch effects associated with laboratory, plate, scanner, platform, or study site. If outcome labels are unevenly distributed across these technical groups, a model can use batch as a surrogate for outcome. Soneson and Delorenzi (2014) and Hamdan, Love, and von dem Hagen (2023) showed that such confounding, and even attempted confound removal, can distort cross-validation-based estimates. In a multi-study setting, however, the intended validation target matters: if one study is held out at a time for external validation, perfect fold-study association is expected and should not be interpreted as a software error. This is one instance of the more general principle that resampling must respect the dependence structure relevant to the intended deployment target (Roberts *et al.* 2017).

2.4. Temporal look-ahead leakage

When data have a temporal order, a random split permits the training set to use information that would not have been available at the prediction time of the test set. The most direct example is an explicit future variable, such as a subsequent laboratory measurement or a later event summary, included as a predictor. A subtler example arises when preprocessing borrows strength across future observations. Proper time-aware resampling therefore requires chronological ordering and, in some settings, purge or embargo windows to prevent overlap between training and evaluation periods. The same point appears in broader cross-validation methodology: evaluation is meaningful only when the data split matches the prediction task that the model is meant to serve (Stone 1974; Arlot and Celisse 2010; Hastie, Tibshirani, and Friedman 2009).

2.5. Target leakage and duplicate samples

Target leakage occurs when a predictor carries direct information about the outcome definition. In electronic health records this may be a billing code or medication that is downstream of the diagnosis. In multi-omics or public expression resources it may be a transformed outcome summary or a feature engineered using labels. Duplicate and near-duplicate samples are related but distinct: they need not encode the outcome explicitly, yet a nearly identical training sample can make a test sample artificially easy to predict. Because such problems may survive apparently careful resampling, they are a natural target for post hoc diagnostics. In audit terms, near-duplicate detection is closer to record linkage or similarity search than to classical hypothesis testing, and it inherits familiar design choices about feature representation, matching rules, and thresholds (Winkler 1999; Christen 2012; Leskovec, Rajaraman, and Ullman 2014).

3. Overview of the `bioLeak` Package

The package is organized around four design goals: leakage-aware modeling, reproducible workflows, diagnostic auditing, and compatibility with common R modeling tools. The software focuses on biomedical prediction studies in which the user already has a labeled dataset and wants to evaluate whether reported performance may depend on leakage or confounding.

Figure 1 summarizes the main package workflow from a labeled dataset to fitted resampling

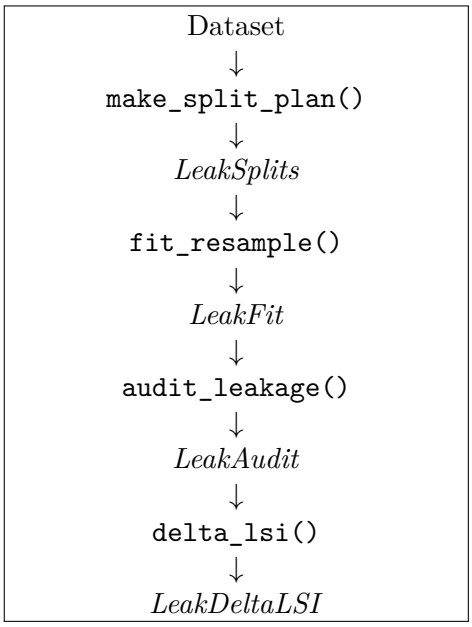

Figure 1: `bioLeak` leakage-aware modeling workflow. Each function produces an S4 object that stores intermediate state for the next stage.

results, audit summaries, and inflation quantification.

Each stage produces an S4 object: `LeakSplits`, `LeakFit`, `LeakAudit`, and `LeakDeltaLSI`. These objects store all intermediate state needed for downstream analysis and reporting.

The main workflow begins with `make_split_plan()`, which constructs an S4 `LeakSplits` object. The split plan records the resampling mode, fold-level indices or compact fold assignments, and metadata such as the grouping variables, number of folds, number of repeats, and a reproducibility hash. The split modes most relevant to the present manuscript are `subject_grouped`, `batch_blocked`, `study_loocv`, and `time_series`. These modes address distinct dependence structures rather than acting as interchangeable cross-validation variants.

Model fitting is handled by `fit_resample()`, which applies guarded preprocessing independently within each fold, trains one or more learners, and returns a `LeakFit` object. The fit object stores per-fold metrics, aggregated summaries, audit metadata, predictions, preprocessing state, learners, and information required for later refit-based audits. The function accepts data frames, matrices, and `SummarizedExperiment` objects. Learners may be specified through `parsnip`, `workflows`, legacy character interfaces, or custom fit/predict pairs.

`tune_resample()` extends this design to nested resampling. Outer folds are defined by a `LeakSplits` object or an `rsample` structure that already contains inner folds. Tuning results are returned as a `LeakTune` object containing inner tuning outputs, outer-fold metrics, selected hyperparameters, optional threshold selections for binary outcomes, and an optional refitted final workflow. Survival tasks are currently excluded from `tune_resample()`.

Auditing is provided by `audit_leakage()`, which returns a `LeakAudit` object. The audit combines a permutation-gap analysis, fold association tests for batch or study metadata, target-leakage scans on a reference feature matrix, duplicate detection, and a compact mechanism-level summary stored in the object metadata. `audit_report()` renders these

Table 1: Primary software components in `bioLeak`.

Stage	Primary function(s)	Main output or role
Split construction	`make_split_plan()`	Leakage-aware split plans for subject, batch, study, or time-ordered resampling.
Cross-validated fitting	`fit_resample()`	Fold predictions, fold metrics, aggregated performance summaries, preprocessing state, and refit metadata.
Nested tuning	`tune_resample()`	Nested cross-validation for hyperparameter tuning with leakage-aware outer and inner folds.
Leakage auditing	`audit_leakage()`	Permutation-gap summaries, batch or study association tests, target scans, duplicate detection, and mechanism summaries.
Audit reporting	`audit_report()`	Self-contained HTML summaries suitable for review and supplementary material.
Inflation quantification	`delta_lsi()`	Repeat-level comparison of naive and guarded pipelines, including robust inflation estimates and randomization-based inference.
Interoperability	`guard_to_recipe()`, `as_rsample()`	Bridges to `recipes` and `rsample` workflows.

results to a self-contained HTML report based on a bundled R Markdown template. For direct inflation comparisons between a naive and a guarded pipeline, the package also provides `delta_lsi()`, which returns a `LeakDeltaLSI` object summarizing repeat-level performance differences, robust point estimates, confidence intervals, and sign-flip randomization tests.

Table 1 lists the primary user-facing functions and their roles in the workflow.

4. Leakage-Aware Modeling Workflow

This section illustrates a typical workflow using a small synthetic binary classification example. The same workflow structure extends to multiclass, regression, and survival tasks, although the available metrics and some audit components differ. The underlying evaluation logic follows standard resampling principles, but it adapts them to structured biomedical data in which the independence unit may be subject, batch, study, or time rather than row (Stone 1974; Kohavi 1995; Arlot and Celisse 2010; James, Witten, Hastie, and Tibshirani 2013).

4.1. Constructing a split plan

The first step is to define the dependence structure that the resampling design should respect. In the example below, repeated measurements are nested within subject, so a subject-grouped split plan is used. The code mirrors the current package API and can be adapted to

`batch_blocked`, `study_loocv`, or `time_series` by replacing the relevant mode and metadata columns.

```
library(bioLeak)
library(parsnip)

set.seed(1)
n_subject <- 40
rep_per_subject <- 3
n <- n_subject * rep_per_subject

subject <- rep(seq_len(n_subject), each = rep_per_subject)
batch <- rep(seq_len(6), length.out = n)
x1 <- rnorm(n)
x2 <- rnorm(n)
x3 <- rnorm(n)
subject_signal <- rnorm(n_subject)
outcome <- factor(
  ifelse(runif(n) < plogis(subject_signal[subject]), "case", "control"),
  levels = c("control", "case")
)

df <- data.frame(subject, batch, outcome, x1, x2, x3)

splits <- make_split_plan(
  df,
  outcome = "outcome",
  mode = "subject_grouped",
  group = "subject",
  v = 5,
  stratify = TRUE,
  seed = 1
)

splits

LeakSplits object (mode = subject_grouped, v = 5, repeats = 1)
Outcome: outcome | Stratified: TRUE | Nested: FALSE
-------------------------------------------------------
  repeat fold n_train n_test
       1    1      96     24
       1    2      96     24
       1    3      96     24
       1    4      96     24
       1    5      96     24
-------------------------------------------------------
Total folds: 5 | Hash: a1b2c3d4e5f6
```

The split object records fold membership and split metadata. For large datasets, `compact = TRUE` stores fold assignments rather than explicit index lists to reduce memory use.

4.2. Defining guarded preprocessing

`fit_resample()` accepts either a native guarded preprocessing specification or a `recipes::recipe`. The native interface applies each step on the training fold only and transfers the resulting parameters to the test fold. Supported operations include winsorization, imputation, normalization, filtering, and feature selection. In the example below, median imputation and z-score normalization are performed within each training fold.

```
spec <- logistic_reg(mode = "classification") |>
  set_engine("glm")

fit_guarded <- fit_resample(
  df,
  outcome = "outcome",
  splits = splits,
  learner = spec,
  metrics = "auc",
  preprocess = list(
    impute = list(method = "median"),
    normalize = list(method = "zscore"),
    filter = list(var_thresh = 0, iqr_thresh = 0),
    fs = list(method = "none")
  ),
  refit = FALSE,
  seed = 1
)

summary(fit_guarded)
```

```
===========================
 bioLeak Model Fit Summary
===========================

Task: binomial
Outcome: outcome
Positive class: case
Learners: logistic_reg
Total folds: 5
Fold status: 5 success, 0 skipped, 0 failed
Refit performed: No
Hash: 7a3f2e1b4c9d

Cross-validated metrics (mean +/- SD):
  metric learner    mean      sd
     auc     glm  0.5417  0.1024
```

For multiclass classification, the same interface can be used with metrics such as `accuracy`, `macro_f1`, and `log_loss`. For survival analysis, the outcome is specified as a two-element character vector containing time and event columns, and the principal metric is the concordance index. For continuous regression, the outcome is a numeric variable, and the available metrics include `rmse`, `mae`, and R^2.

4.3. Training naive and guarded comparators

The package is most informative when a guarded analysis is accompanied by a plausible naive or intentionally leaky comparator. The next snippet adds a subject-level leakage proxy and fits an otherwise identical pipeline. Such comparisons are useful both for exploratory sensitivity analysis and for the `delta_lsi()` procedure described in Section 5.

```
df_leaky <- within(df, {
  leak_subject <- ave(as.numeric(outcome == "case"), subject, FUN = mean)
})

fit_naive <- fit_resample(
  df_leaky,
  outcome = "outcome",
  splits = splits,
  learner = spec,
  metrics = "auc",
  preprocess = list(
    impute = list(method = "none"),
    normalize = list(method = "none"),
    filter = list(var_thresh = 0, iqr_thresh = 0),
    fs = list(method = "none")
  ),
  refit = FALSE,
  seed = 1
)
```

4.4. Auditing a fitted workflow

After cross-validated predictions have been generated, `audit_leakage()` can be used to summarize the evidence available within the fitted object and, optionally, in an aligned reference matrix `X_ref`. The reference matrix enables target-leakage scans and duplicate detection. In grouped designs, the same audit interface can be used with restricted permutations and batch-association tests.

```
audit_guarded <- audit_leakage(
  fit_guarded,
  metric = "auc",
  B = 50,
  perm_stratify = TRUE,
```

```
  X_ref = df[, c("x1", "x2", "x3")],
  batch_cols = "batch",
  return_perm = TRUE
)

summary(audit_guarded)
```

```
==============================
 bioLeak Leakage Audit Summary
==============================

Task: binomial | Outcome: outcome | Splitting mode: subject_grouped
Hash: 7a3f2e1b4c9d | Folds: 5 | Repeats: 1

Label-Permutation Association Test:
  Method: fixed predictions
  Observed metric: 0.5417
  Permuted mean +/- SD: 0.5012 +/- 0.0389
  Gap: 0.0405 (larger gap = stronger non-random signal)

Batch / Study Association:
  batch (repeat 1): Chi^2 = 4.123 (df = 4), p = 0.3897

Target Leakage Scan:
  Features checked: 3 | Flagged (score >= 0.90): 0

Multivariate Target Scan: not available.

Near-Duplicate Samples:
  No near-duplicates detected.

Mechanism Risk Assessment:
  mechanism_class      flagged evidence
  subject_overlap      FALSE   OK: subject-grouped splits
  batch_confounded     FALSE   OK: batch Chi^2 not significant
  preprocessing_leak   FALSE   OK: guarded preprocessing
  target_leakage       FALSE   OK: no features flagged

Interpretation:
  Modest non-random signal. No leakage indicators flagged.
```

The resulting `LeakAudit` object stores one-row permutation summaries, feature-level target-association results, duplicate pairs, and a trail of audit parameters. The permutation module is diagnostic rather than dispositive: a large positive gap indicates non-random predictive signal, but the other audit components are needed to assess whether that signal may be inflated by leakage.

4.5. Generating an HTML audit report

The audit can be rendered into an HTML report without modifying the underlying fit object. The report is particularly useful when a model review requires a human-readable appendix.

```
report_file <- audit_report(
  audit_guarded,
  output_file = "bioLeak_audit_report.html",
  output_dir = tempdir(),
  quiet = TRUE
)
```

4.6. Putting the workflow together

The previous subsections present each stage separately. In practice, the full workflow is usually executed as one analysis script that constructs a split plan, fits guarded and naive comparators, audits the guarded fit, quantifies inflation, and renders a report. The following code block illustrates that sequence in one place.

```
library(bioLeak)
library(parsnip)

set.seed(1)
n_subject <- 30
rep_per_subject <- 4
n <- n_subject * rep_per_subject

subject <- rep(seq_len(n_subject), each = rep_per_subject)
batch <- rep(seq_len(6), length.out = n)
latent <- rnorm(n_subject)
x1 <- latent[subject] + rnorm(n, sd = 0.5)
x2 <- rnorm(n)
x3 <- rnorm(n)
outcome <- factor(
  ifelse(runif(n) < plogis(latent[subject]), "case", "control"),
  levels = c("control", "case")
)

dat <- data.frame(subject, batch, outcome, x1, x2, x3)

plan <- make_split_plan(
  dat,
  outcome = "outcome",
  mode = "subject_grouped",
  group = "subject",
  v = 5,
  stratify = TRUE,
```

```
  seed = 1
)

spec <- logistic_reg(mode = "classification") |>
  set_engine("glm")

fit_guarded <- fit_resample(
  dat,
  outcome = "outcome",
  splits = plan,
  learner = spec,
  metrics = "auc",
  preprocess = list(
    impute = list(method = "median"),
    normalize = list(method = "zscore"),
    filter = list(var_thresh = 0, iqr_thresh = 0),
    fs = list(method = "none")
  ),
  refit = FALSE,
  seed = 1
)

dat_naive <- within(dat, {
  leak_subject <- ave(as.numeric(outcome == "case"), subject, FUN = mean)
})

fit_naive <- fit_resample(
  dat_naive,
  outcome = "outcome",
  splits = plan,
  learner = spec,
  metrics = "auc",
  preprocess = list(
    impute = list(method = "none"),
    normalize = list(method = "none"),
    filter = list(var_thresh = 0, iqr_thresh = 0),
    fs = list(method = "none")
  ),
  refit = FALSE,
  seed = 1
)

audit <- audit_leakage(
  fit_guarded,
  metric = "auc",
  B = 50,
  perm_stratify = TRUE,
```

```
  batch_cols = "batch",
  X_ref = dat[, c("x1", "x2", "x3")]
)

delta <- delta_lsi(
  fit_naive = fit_naive,
  fit_guarded = fit_guarded,
  metric = "auc",
  seed = 1
)

report_file <- audit_report(
  audit,
  output_file = "bioLeak_audit_report.html",
  output_dir = tempdir(),
  quiet = TRUE
)
```

The split plan defines the dependence structure, `fit_resample()` trains the guarded and naive pipelines under that design, `audit` stores diagnostic evidence for the guarded fit, `delta` summarizes the repeat-level inflation between the naive and guarded analyses, and `audit_report()` produces a shareable HTML summary. The same pattern extends to multiclass, regression, and survival settings with task-appropriate learners and metrics.

4.7. Leakage-safe hyperparameter tuning

`tune_resample()` performs nested cross-validation so that model selection remains inside the outer resampling loop. The example below uses a tunable `glmnet` specification and an inner resampling structure derived from a nested split plan. This design avoids tuning on the same assessment data used for outer-loop evaluation.

```
library(recipes)
library(tune)

nested_splits <- make_split_plan(
  df,
  outcome = "outcome",
  mode = "subject_grouped",
  group = "subject",
  v = 3,
  nested = TRUE,
  seed = 14
)

rec <- recipe(outcome ~ x1 + x2 + x3, data = df) |>
  step_impute_median(all_numeric_predictors()) |>
  step_normalize(all_numeric_predictors())
```

```
spec_tune <- logistic_reg(
  penalty = tune(),
  mixture = 1,
  mode = "classification"
) |>
  set_engine("glmnet")

tuned <- tune_resample(
  df,
  outcome = "outcome",
  splits = nested_splits,
  learner = spec_tune,
  preprocess = rec,
  grid = 5,
  metrics = c("auc", "accuracy"),
  selection = "one_std_err",
  refit = TRUE,
  seed = 14
)
```

When a `recipe` or `workflow` is used, `bioLeak` checks the graph for operations that may violate fold-wise estimation rules. The package also provides `guard_to_recipe()` when a guarded preprocessing specification needs to be translated into an approximate `recipes` pipeline.

4.8. Practical workflow variants

The same structure generalizes to the other supported split modes. With `batch_blocked`, all records from the same batch are isolated in one fold. With `study_loocv`, one study is held out at a time, which is often appropriate for consortium data or data aggregated from public resources. With `time_series`, observations are ordered by time and training is restricted to records that precede the test block, optionally with a prediction horizon, purge, or embargo window. In all cases, the core principle remains the same: resampling should represent the target deployment setting rather than a convenient random partition of rows.

5. Quantifying Leakage Inflation (`Delta LSI`)

`bioLeak` provides `delta_lsi()` to compare a naive pipeline with a guarded one when both are fitted on the same repeated resampling design. The goal is not to decide whether one pipeline is "better" in a generic sense, but to quantify how much apparent performance changes when leakage controls are introduced. This focus is consistent with the broader cross-validation literature, which emphasizes that resampling estimates are themselves random quantities and that their uncertainty is nontrivial to characterize (Bengio and Grandvalet 2004; Bates, Hastie, and Tibshirani 2024).

Suppose repeated cross-validation yields R repeat-level summaries. Let $M_{\text{naive},r}$ and $M_{\text{guarded},r}$ denote the repeat-level performance estimates for repeat r. For higher-is-better metrics such

as AUC or accuracy, the repeat-level inflation score is

$$\Delta_r = M_{\text{naive},r} - M_{\text{guarded},r}, \qquad r = 1, \ldots, R.$$

Positive values therefore indicate that the naive pipeline appears more optimistic than the guarded one and, under the paired design, are interpreted as leakage-related performance inflation. For lower-is-better metrics such as RMSE or log-loss, `delta_lsi()` applies a sign flip internally through the `higher_is_better` argument so that positive values still denote inflation.

In the implementation, fold-level metric values are aggregated to a repeat level using test-set-size weights. Let μ_r denote the resulting repeat summary for one pipeline. The vector $\boldsymbol{\Delta} = (\Delta_1, \ldots, \Delta_R)^\top$ is then summarized in two ways. The first is the arithmetic mean,

$$\widehat{\Delta}_{\text{metric}} = \frac{1}{R} \sum_{r=1}^{R} \Delta_r,$$

which is easy to interpret directly in the units of the metric. The second is the robust summary reported as `delta_lsi`,

$$\widehat{\Delta}_{\text{LSI}} = \arg\min_{\mu} \sum_{r=1}^{R} \rho_{1.345} \left(\frac{\Delta_r - \mu}{\widehat{\sigma}} \right),$$

where $\rho_{1.345}$ is Huber's loss function and $\widehat{\sigma} = 1.4826 \times \text{MAD}(\boldsymbol{\Delta})$ is a fixed robust scale estimate. The Huber estimator downweights outlying repeats, which is useful when one or two resampling repeats are unstable for reasons unrelated to systematic leakage (Huber 1964; Huber and Ronchetti 2009).

Confidence intervals are computed with a bias-corrected and accelerated (BCa) bootstrap when the effective number of paired repeats is at least ten. The package reports a 95% interval for both `delta_lsi` and the arithmetic mean. Hypothesis testing is based on a sign-flip randomization test. Under the null hypothesis of no systematic inflation, the sign of each Δ_r is exchangeable. For $R \leq 15$, the package enumerates all 2^R sign combinations exactly. For larger R, Monte Carlo sign-flips are sampled and the Phipson-Smyth correction is applied so that the reported p-value is never zero. When `exchangeability = "blocked_time"`, a block sign-flip procedure is used instead of the independent sign-flip test. In this sense, the reported significance summary is a permutation-style randomization test on the repeat-level differences rather than a model-based large-sample approximation. The use of bootstrap intervals follows standard resampling-based inference ideas (Efron 1979; Efron and Tibshirani 1993).

The design is only fully inferential when the two compared fits are paired. In practice this means they must share the same fold structure, not merely the same number of repeats. If fold membership differs, the package falls back to an unpaired comparison of repeat-level summaries and suppresses formal inference. This distinction is important. If the user changes both the predictors and the resampling design, the resulting performance difference cannot be attributed only to leakage-aware preprocessing.

The returned `LeakDeltaLSI` object stores the metric name, the exchangeability assumption, the effective number of paired repeats, robust and arithmetic point estimates, bootstrap intervals, the sign-flip p-value, fold- and repeat-level summaries for both pipelines, and an `info` slot with pairing diagnostics and optional details. The package groups inferential output into four tiers:

1. `D_insufficient`: fewer than five effective paired repeats or an unpaired comparison; point estimates only.
2. `C_signflip`: at least five paired repeats; point estimates and a sign-flip p-value.
3. `B_signflip_ci`: at least ten paired repeats; point estimates, p-value, and BCa intervals.
4. `A_full_inference`: at least twenty paired repeats with finite intervals and p-value.

A typical call compares a naive and a guarded fit built from the same `LeakSplits` object:

```
inflation <- delta_lsi(
  fit_naive = fit_naive,
  fit_guarded = fit_guarded,
  metric = "auc",
  M_boot = 2000,
  M_flip = 10000,
  seed = 42
)

summary(inflation)
```

```
====================================
 bioLeak Delta LSI Summary
====================================

Metric:            auc
Exchangeability:   iid
Inference tier:    A_full_inference
R_eff:             20  (R_naive=20, R_guarded=20, paired=TRUE)

Naive pipeline:    0.720
Guarded pipeline:  0.612

Point estimates (naive - guarded; positive = leakage inflation):
  delta_metric:  0.108  (raw metric difference)
  95% BCa CI:    [0.061, 0.155]
  delta_lsi:     0.107  (Huber-robust)
  95% BCa CI:    [0.060, 0.154]

Hypothesis test  [H0: mean(dr) = 0; sign-flip tests delta_metric]:
  Sign-flip p:   0.0001 ***

Inference valid:   YES (tier A)
```

In practice, `delta_lsi()` is most useful when an analyst wants to measure the sensitivity of a result to specific leakage controls, such as subject grouping, batch blocking, or removal of explicit leak-prone features. The case study in Section 7 demonstrates a paired comparison,

and Appendix B reports a small simulation evaluating the Type I error and power of the sign-flip test.

6. Simulation Study

6.1. Design

An important interpretive caveat applies throughout this section. All permutation-gap analyses use `perm_refit = FALSE`, meaning that the model is fitted once on the original labels and the permutation test shuffles labels against the fixed predictions. This design tests whether the prediction-label association exceeds chance levels, but it does not distinguish between genuine predictive signal and signal attributable to leakage. As a consequence, the clean baseline (`none`) also rejects at high rates when real signal is present ($s > 0$). The leakage-specific detection power is therefore best assessed at $s = 0$, where no real signal exists and the only source of above-chance performance is the injected leakage mechanism.

The paper repository contains a dedicated simulation script, `paper/run_simulation.R`, whose design is intentionally distinct from the package-level helper `simulate_leakage_suite()`. The manuscript results therefore describe the stored paper artifacts rather than package defaults. The main differences are: (i) the paper generates `peek_norm` leakage as $y + \varepsilon$ with $\varepsilon \sim \mathcal{N}(0, 0.09)$, whereas the package function uses a z-scored binary outcome; (ii) the paper generates `lookahead` leakage by shifting a continuous biomarker rather than the binary outcome; and (iii) the paper adds AR(1) noise ($\rho = 0.9$) to the latent linear predictor when $s > 0$. These choices were made to produce leakage signals of graded strength; readers who wish to reproduce the manuscript results exactly should use the paper script, not the package helper. The main result object, `paper/sim_results_all.rds`, contains 3,600 rows corresponding to a full factorial design with

- five leakage settings: `none`, `subject_overlap`, `batch_confounded`, `peek_norm`, and `lookahead`;
- four sample sizes: $n \in \{100, 250, 500, 1000\}$;
- three feature dimensions: $p \in \{10, 50, 100\}$;
- three signal levels: $s \in \{0, 0.5, 2.0\}$; and
- twenty seeds per configuration.

For each simulated dataset, predictors were generated as Gaussian variables and the binary outcome was sampled from a probit link applied to a latent linear predictor. When $s = 0$, the outcome was pure noise and served as a type-I error baseline. When $s > 0$, signal entered through the first few predictors, and autoregressive noise was added to the latent process. Leakage was injected as follows.

- `subject_overlap`: a subject-level mean of the binary outcome was added as a predictor.
- `batch_confounded`: batch membership was made outcome dependent and the batch-wise mean outcome was added as a predictor.

- `peek_norm`: a noisy continuous version of the outcome was added as a globally computed feature.
- `lookahead`: a future biomarker value was shifted into the current row.

All main-simulation analyses used five-fold subject-grouped cross-validation with `glmnet` as the learner, median imputation, z-score normalization, no feature selection, and a permutation-gap audit with $B = 50$ and `perm_refit = FALSE`. The stored columns include the observed AUC, permutation-null mean, gap, and permutation p-value.

Two supplementary result objects refine this design. The file `paper/sim_results/supplementary_modes.rds` evaluates the same general leakage mechanisms under four split modes (`subject_grouped`, `batch_blocked`, `study_loocv`, and `time_series`) at a fixed setting ($n = 500, p = 20, s = 1.0$). The file `paper/sim_results/supplementary_target_scan.rds` evaluates the target-scanning components at the same fixed setting. Scalar summaries derived from these objects were also written to `paper/placeholder_values.rds` via `paper/compile_results.R`.

6.2. Main simulation summaries

```
sim <- readRDS("paper/sim_results_all.rds")
head(sim)
```

```
  seed   n   p   s     leakage metric_obs perm_mean    gap p_value
1    1 100  10 0.0        none     0.5200    0.5014  0.019    0.24
2    2 100  10 0.0        none     0.4850    0.4998 -0.015    0.64
...
```

The clean baseline in `sim_results_all.rds` provides a direct check of the permutation-gap procedure under $s = 0$ and no injected leakage. Across all feature dimensions and sample sizes, the proportion of baseline runs with $p < 0.05$ was 6.7%. Broken down by sample size, the corresponding values were 10.0% for $n = 100$, 8.3% for $n = 250$, 1.7% for $n = 500$, and 6.7% for $n = 1000$. These figures are slightly variable around the nominal level, as expected with only twenty seeds per configuration, but they do not suggest a systematic failure of the stored null runs.

For leakage-injected settings with $s > 0$, the permutation-gap statistic was usually significant. However, with `perm_refit = FALSE` the permutation-gap test detects prediction-label association rather than leakage specifically. Crucially, the clean baseline (`none`) also rejects at high rates when real signal is present: 84.4% overall at $s > 0$, reaching 99% at $n = 500$ and 100% at $n = 1000$. Table 2 includes this baseline alongside the leakage-injected runs so that the increment attributable to leakage is visible. Pooled over feature dimension and the two positive signal levels, the weakest leaky setting is `lookahead` at $n = 100$ (0.62), which is close to the clean baseline at the same n (0.56). All other mechanisms exceed the clean baseline by a larger margin at $n = 100$ and converge to 1.00 by $n = 500$.

Entries are based on 20 seeds per cell. For a proportion $\hat{p}$, the standard error is approximately $\sqrt{\hat{p}(1-\hat{p})/20}$, yielding ± 0.11 at $\hat{p} = 0.50$ and ± 0.04 at $\hat{p} = 0.90$.

These values should be interpreted carefully. Because the permutation-gap test with `perm_refit = FALSE` tests non-random predictive signal rather than leakage in isolation, the high baseline

Table 2: Proportion of simulation runs with permutation-gap $p < 0.05$, pooled over feature dimension and signal level $s > 0$, based on `paper/sim_results_all.rds`. The first row (None) is the clean baseline; the increment above this row is attributable to leakage.

Mechanism	$n = 100$	$n = 250$	$n = 500$	$n = 1000$
None (clean baseline)	0.56	0.82	0.99	1.00
Subject overlap	0.90	1.00	1.00	1.00
Batch confounded	0.95	1.00	1.00	1.00
Peek normalization	1.00	1.00	1.00	1.00
Look-ahead	0.62	0.94	1.00	1.00

Table 3: Leakage-specific detection at $s = 0$ (no real signal). Entries are proportions of runs with permutation-gap $p < 0.05$, pooled over feature dimension. The clean baseline (`none`) is 6.7%.

Mechanism	$n = 100$	$n = 250$	$n = 500$	$n = 1000$
Subject overlap	0.78	1.00	1.00	1.00
Batch confounded	0.88	1.00	1.00	1.00
Peek normalization	1.00	1.00	1.00	1.00
Look-ahead	0.12	0.08	0.05	0.05

rejection rate at $s > 0$ means that absolute rejection rates cannot be read as leakage-detection power. A more informative view comes from the $s = 0$ condition, where no real signal exists and the only source of above-chance performance is the injected leakage. Table 3 shows the leakage-specific detection rates under this design. Three of the four mechanisms are detected reliably by $n = 250$; `lookahead` is not detectable at $s = 0$ because the shifted biomarker carries no outcome information when the latent signal is absent.

Each cell is based on 20 seeds. Standard errors follow the same binomial formula as Table 2.

This two-table presentation separates the permutation test's role as a general performance check ($s > 0$) from its role as a leakage detector ($s = 0$). The package combines the permutation module with targeted diagnostics for confounding, duplicates, and target leakage precisely because no single component is leakage-specific in all settings.

To quantify inflation directly, the paper workflow compares each leakage mechanism with the clean baseline at matched (n, s) values. Averaged over the stored configurations with $s > 0$, the mean observed AUC in clean settings was 0.75. Relative to that clean baseline, the mean inflation in observed AUC was 0.086 for `subject_overlap`, 0.091 for `batch_confounded`, 0.243 for `peek_norm`, and 0.026 for `lookahead`. In these stored results, explicit global peek-type leakage produces the largest inflation, whereas the look-ahead construction yields the smallest average AUC increase. This difference is attributable to the particular manuscript design in `paper/run_simulation.R`; it should not be read as a universal ranking of leakage severity. Figure 2 summarizes these patterns graphically.

6.3. Supplementary split-mode and target-scan analyses

Table 4 summarizes the supplementary split-mode object `paper/sim_results/supplementary_modes.rds`. As with the main simulation, the split-mode experiment was run at both $s = 0$ (leakage-

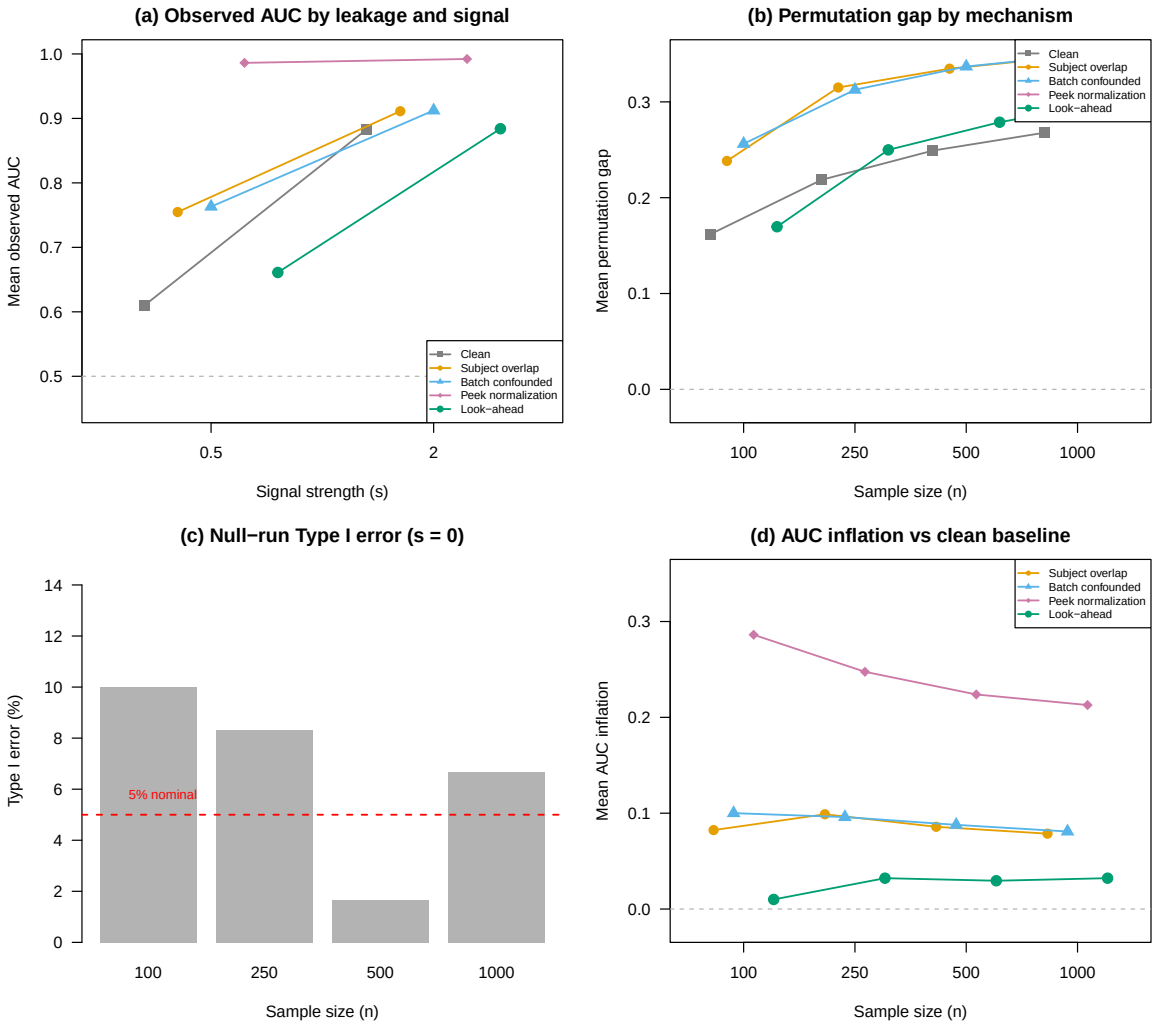

Figure 2: Main simulation summaries from `paper/sim_results_all.rds`. (a) Observed AUC by leakage mechanism and signal level. (b) Permutation-gap summaries by mechanism and sample size. (c) Null-run type I error at $s = 0$. (d) AUC inflation relative to the clean baseline.

Table 4: Supplementary split-mode analysis based on `paper/sim_results/supplementary_modes.rds`. Entries are the proportions of leakage-injected runs with permutation-gap $p < 0.05$ at $(n = 500, p = 20, s = 0)$. The clean baseline at $s = 0$ rejects at approximately 5% for all modes.

Split mode	**Subject overlap**	**Batch confounded**	**Peek normalization**	**Look-ahead**
Subject grouped	1.00	1.00	1.00	0.05
Batch blocked	1.00	0.00	1.00	0.00
Study LOOCV	1.00	1.00	1.00	0.00
Time series	1.00	1.00	1.00	0.00

specific) and $s = 1.0$ (signal present). At $s = 1.0$, the clean baseline rejects in all 20/20 runs for every split mode, so the $s = 1.0$ results cannot isolate leakage. Table 4 therefore reports the $s = 0$ condition, where the only source of above-chance performance is the injected leakage. Under all four split modes, `subject_overlap`, `batch_confounded`, and `peek_norm` are detected with high power, with the notable exception of `batch_confounded` under `batch_blocked` mode (0.00), where the blocking variable aligns with the confounding structure. Look-ahead leakage is largely undetectable at $s = 0$ (rejection rates 0.00–0.05 across modes), which is expected: this mechanism leaks future values of a continuous biomarker that carries signal only when real signal is present.

The reduced detection of `batch_confounded` leakage under `batch_blocked` mode is an expected interaction: when the test fold contains samples from a single batch, confounding between batch and outcome is inherently difficult to detect via permutation because the batch structure is held constant. This illustrates a key design lesson: the permutation-gap test's sensitivity depends on the interplay between the leakage mechanism and the resampling geometry. The exercise reinforces the main software point of `bioLeak`: the choice of split design is part of the model specification, and combining permutation tests with targeted diagnostics (such as the target scan) provides more robust detection.

The supplementary target-scan object, `paper/sim_results/supplementary_target_scan.rds`, provides a more granular picture. At $s = 0$ (no real signal), the univariate target scan flagged the injected leakage feature in all 20/20 `peek_norm` runs and in none of the `subject_overlap`, `batch_confounded`, `lookahead`, or clean runs at the threshold used in the stored script (`target_threshold = 0.9`). Averaged over all leaky conditions, this corresponds to a 25% success rate because only one of the four injected mechanisms satisfied the conservative univariate threshold. The multivariate scan at $s = 0$ rejected in 76% of leaky runs, with a 5% false-positive rate in the clean baseline—consistent with the nominal $\alpha = 0.05$ level. This contrasts with the $s = 1.0$ evaluation, where the multivariate scan rejects in all conditions (including the clean control) because the real signal features dominate the PCA–GLM statistic. The $s = 0$ results confirm that the multivariate scan has genuine leakage-detection power when confounding with real signal is removed. Nonetheless, the univariate and multivariate scans should be interpreted jointly, and with knowledge of the reference matrix supplied to `X_ref`. Figure 3 visualizes these results.

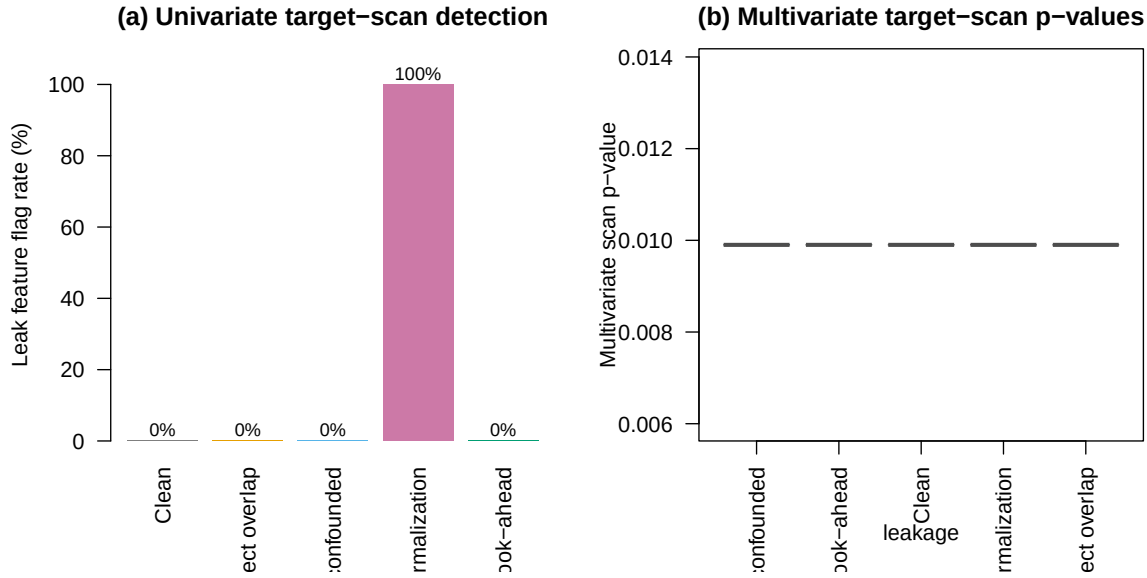

Figure 3: Supplementary target-scan summaries from `paper/sim_results/supplementary_target_scan.rds`. (a) Univariate leak-feature flag rate by mechanism. (b) Multivariate target-scan p-values.

7. Case Study

7.1. Data and workflow

The case study uses the stored object `paper/casestudy_results.rds`, which was generated by `paper/run_casestudy.R`. The script draws data from the `curatedOvarianData` Bioconductor package and constructs a binary three-year overall survival endpoint from studies containing both `days_to_death` and `vital_status`. After intersecting feature sets across eligible studies and retaining the 2,000 most variable common genes, the final analysis dataset contains $N = 2015$ samples from $S = 14$ studies and $G = 2000$ genes. Study sizes range from 50 to 376 samples, with a median of 125. Multi-study gene-expression prediction has long served as a test case for both high-dimensional classifiers and evaluation pitfalls in biomedical machine learning (Golub, Slonim, Tamayo, Huard, Gaasenbeek, Mesirov, Coller, Loh, Downing, Caligiuri, Bloomfield, and Lander 1999; Simon *et al.* 2003; Tarca, Carey, Chen, Romero, and Drăghici 2007).

```
cs <- readRDS("paper/casestudy_results.rds")
names(cs)
```

```
[1] "N"         "S"         "G"         "guarded"   "leaky"     "delta_lsi"
```

Two workflows were evaluated. The guarded workflow used `study_loocv`, median imputation, z-score normalization, t-test feature selection with the top 100 genes per fold, and a `glmnet` learner. The naive comparator used standard five-fold row-wise evaluation implemented via a `subject_grouped` split on a dummy row identifier, the same learner, and three deliberately

Table 5: Stored case-study summary from `paper/casestudy_results.rds`.

Workflow	**Split design**	**AUC**	**SD**	**Gap**	**p-value**	**Study-fold V**	**Multivariate sc**
Guarded	Study LOOCV	0.612	0.047	0.053	0.0020	1.000	0.0792
Naive/leaky	Row-wise 5-fold CV	0.720	0.026	0.211	0.0020	0.079	0.0099

leak-prone features: a study-level outcome mean, a globally standardized outcome value, and a globally standardized first principal component.

The two workflows differ not only in their leakage controls but also in preprocessing: the guarded pipeline applies per-fold t-test feature selection (top 100 genes), whereas the naive comparator uses all available features. This asymmetry means that the observed performance difference reflects both the removal of leakage and the effect of dimensionality reduction. The design was chosen to represent a realistic contrast between a carefully constructed and a carelessly constructed pipeline, but the AUC gap should not be attributed solely to leakage removal.

Both workflows were audited with a permutation-gap test ($B = 500$), study-association testing, target scans, and duplicate detection.

7.2. Stored results

Table 5 summarizes the stored case-study output. The guarded workflow achieved an AUC of 0.612 with fold-to-fold standard deviation 0.047, whereas the naive comparator achieved an AUC of 0.720 with standard deviation 0.026. The difference in observed AUC is therefore 0.108 in favor of the naive workflow. The permutation-gap component is positive for both workflows, but the gap is much larger in the naive comparator (0.211 versus 0.053).

The study-fold association results deserve explicit interpretation. In the guarded workflow, Cramer's V is exactly 1.000 with a vanishingly small p-value because each outer fold is defined by study. This is not evidence of an implementation error; it is a property of the intended external-validation design. In the naive comparator, the association is weak ($V = 0.079$, $p = 0.5302$) because the row-wise split ignores study boundaries. This contrast illustrates why association tables must be interpreted against the chosen resampling policy rather than in isolation.

The target-scan results are more directly diagnostic. In the guarded workflow, no feature exceeded the conservative univariate threshold. The top stored association scores correspond to genes `XRCC4`, `XPA`, and `XPO1`, with scores 0.167, 0.143, and 0.097, respectively. In the naive comparator, the feature `leak_global_y` was flagged with a score of exactly 1.000, whereas the weaker proxy features `leak_study_mean` and `leak_global_pc1` had scores 0.205 and 0.157 and were not flagged at the stored threshold. The multivariate target scan also distinguished the two workflows: the guarded analysis yielded a stored p-value of 0.0792 (not significant at $\alpha = 0.05$), whereas the naive comparator yielded 0.0099.

Duplicate detection returned a large number of near-identical cross-fold pairs in both workflows: 21,435 pairs in the guarded workflow and 35,220 in the naive workflow. In the stored audit, these pairs were defined by cosine similarity on the reference expression matrix with a reporting threshold of 0.995, retaining pairs that crossed training and test partitions in at least one repeat. The stored manuscript compilation describes these pairs as technical

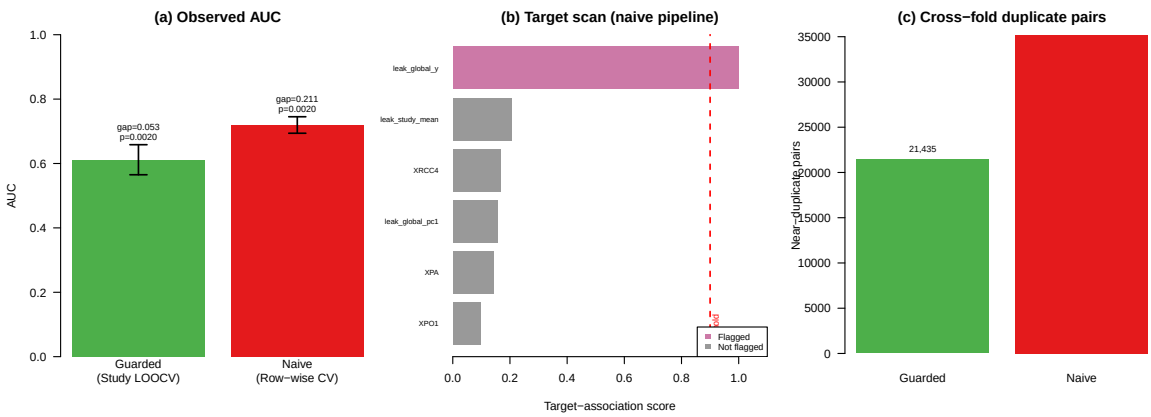

Figure 4: Case-study summaries from `paper/casestudy_results.rds`. (a) Observed AUC with fold standard deviations. (b) Top univariate target-association scores from the naive pipeline. (c) Near-duplicate cross-fold pair counts.

replicates or near-identical expression profiles spanning studies. This result is important for interpretation. It suggests that leakage-aware study holdout alone does not guarantee independence when public data resources contain replicated or near-replicated samples across studies. The duplicate screen therefore contributes information that is not captured by the resampling policy itself. Figure 4 presents these summaries graphically.

To quantify the AUC difference more formally, we applied `delta_lsi()` to the paired fits (both evaluated on the same batch-blocked fold structure with 20 repeats, as documented in `paper/run_delta_lsi.R`). The arithmetic mean inflation was $\widehat{\Delta}_{\text{metric}} = 0.085$ AUC, and the robust Huber estimate was $\widehat{\Delta}_{\text{LSI}} = 0.085$ AUC, with a 95% BCa bootstrap interval of [0.078, 0.091]. The sign-flip randomization test yielded $p < 0.001$, supporting the conclusion that the naive pipeline exhibits systematic performance inflation relative to the guarded one. This pairing is important: because the naive and guarded pipelines share the same fold assignments, the repeat-level differences isolate the effect of leakage controls from the variance of fold sampling. Full `summary()` output for the `LeakDeltaLSI` object is reproduced in Appendix A.

Overall, the case study demonstrates the intended use of the package. A guarded workflow produces a more conservative cross-study estimate, while the audit components identify explicit target leakage and extensive near-duplicate overlap in the naive comparator. The same exercise also shows that some audit components, such as study-fold association, depend on the evaluation target and must be interpreted in context.

8. Software Design and Implementation

8.1. Object structure

The package uses S4 classes to maintain explicit state across the modeling and auditing workflow. `LeakSplits` stores the split mode, fold descriptors, and metadata. `LeakFit` stores the splits, per-fold metrics, metric summaries, predictions, preprocessing state, learners, task

type, and additional metadata such as sample identifiers and refit information. `LeakAudit` stores the audited fit together with permutation summaries, duplicate pairs, target-association tables, batch-association results, and audit metadata. `LeakDeltaLSI` stores paired or unpaired inflation comparisons between two fits.

This separation is useful for reproducibility. Because fits and audits are explicit objects, downstream summaries do not need to recompute the full model pipeline. It also enables auditing and reporting to be decoupled from fitting, which is advantageous when model fitting is computationally expensive.

8.2. Pipeline components

Guarded preprocessing is implemented through an internal `GuardFit` object created by `.guard_fit()`. The training fold is used to estimate all transformation parameters, and a transform function is stored for applying those parameters to the test fold. Supported imputation methods include median imputation, k-nearest-neighbor imputation, and random-forest-based imputation (via optional dependencies). Normalization supports z-score and robust scaling. Filtering supports variance and interquartile-range thresholds. Feature selection includes several univariate and model-based methods, including `ttest`, `lasso`, and `pca`. Character and factor handling is normalized internally so that training-derived levels are applied consistently to the test fold.

Split construction similarly encodes design decisions explicitly. The `time_series` mode orders observations by a supplied time variable and can enforce a prediction horizon, purge window, and embargo window. The `subject_grouped`, `batch_blocked`, and `study_loocv` modes use grouping columns rather than row identities. A compact storage mode reduces memory requirements for large analyses by storing fold assignments rather than explicit train/test vectors.

8.3. Reproducibility features

Several implementation details support reproducible analysis. Split objects store seeds and a content hash. Fit objects can retain sufficient information for refit-based permutation tests through `store_refit_data = TRUE`. Audit objects store both the parameters used and, optionally, the permutation draws themselves. The paper directory in the repository illustrates a workflow in which simulation and case-study scripts write result objects to disk and a separate compilation script extracts the scalar values used in the manuscript. This separation is consistent with a reproducible-paper workflow: expensive computations are executed once, while the paper draws from immutable result artifacts. The same principle is central to the broader reproducible-research literature, which argues that computational claims should remain inspectable at the level of code, data products, and derived summaries (Gentleman and Lang 2007; Peng 2011; Sandve, Nekrutenko, Taylor, and Hovig 2013).

8.4. Reproducing the paper artifacts

The paper directory is organized so that readers can reproduce the manuscript results from stored artifacts without rerunning every computation. The main simulation results are stored in `paper/sim_results_all.rds`, the case study is stored in `paper/casestudy_results.rds`, and the supplementary analyses are stored in `paper/sim_results/supplementary_target_scan.rds`

and `paper/sim_results/supplementary_modes.rds`. The file `paper/placeholder_values.rds` contains scalar summaries used for tables and figure annotations. The scripts `paper/run_simulation.R`, `paper/run_supplementary.R`, `paper/run_casestudy.R`, and `paper/compile_results.R` document how these stored outputs were generated and summarized. A reader who wants to reproduce the reported tables can therefore begin by loading the stored `.rds` files directly and verifying the summaries before rerunning the computationally heavier scripts.

8.5. Dependency structure

The core package depends only on a limited set of packages, but optional dependencies extend functionality. `parsnip` and `hardhat` support modern model specifications. `recipes`, `workflows`, `yardstick`, `tune`, and `dials` support tidymodels interoperability. `glmnet`, `ranger`, and `survival` provide commonly used learners and metrics. `RANN` and `FNN` can accelerate nearest-neighbor duplicate searches. `rmarkdown` is used only when HTML audit reports are rendered. The package therefore remains usable for basic workflows even when some optional tools are not installed.

8.6. Relationship to existing software

`bioLeak` is not a general replacement for R machine-learning frameworks such as `caret`, `tidymodels`, or `mlr3`. Its role is narrower and complementary. Those frameworks provide broad model interfaces, preprocessing tools, and resampling utilities. `bioLeak` focuses specifically on leakage-aware split construction, guarded preprocessing, post hoc leakage diagnostics, and structured reporting. The package can be used independently, but it is designed to coexist with existing modeling tools rather than to supersede them (Kuhn 2008; Kuhn and Wickham 2020; Lang *et al.* 2019; Bischl *et al.* 2024).

9. Discussion

9.1. Software perspective

`bioLeak` addresses a practical gap in biomedical machine-learning workflows by pairing leakage-aware prevention mechanisms with post hoc diagnostic tools. The package design reflects a simple principle: performance estimation and leakage assessment should be treated as parts of one workflow, not as separate stages owned by different pieces of software. This integration is useful in biomedical settings because the causes of leakage are often spread across data preparation, resampling, and study design.

The package has several strengths. First, it makes the resampling geometry explicit. Users must choose whether the relevant independence unit is subject, batch, study, or time. Second, train-fold-only preprocessing is built into the central fitting interface. Third, the auditing framework creates a record of what was checked and how. Fourth, `delta_lsi()` offers a direct way to quantify how much reported performance changes when leakage controls are introduced. Taken together, these components support a style of model assessment that is more transparent than a single cross-validated AUC or accuracy estimate. In that sense, the package makes one practical response to the reproducibility concerns emphasized by Kapoor and Narayanan (2023): leakage checks become part of the analysis record rather than an

informal afterthought.

9.2. Limitations

Several limitations should be stated clearly. First, the package does not prove the absence of leakage. Audit components only test the evidence available in the stored predictions, metadata, and optional reference matrix. If a relevant grouping variable is omitted, no audit procedure can reconstruct it. Second, unsupervised learning workflows are not currently supported; the package is organized around supervised outcomes and supervised or outcome-linked audit procedures. Third, complex graph-based dependency modeling is not yet implemented. Current split construction is based on subject, batch, study, time, and related grouped constraints rather than on arbitrary dependency graphs.

The modeling interface supports binary classification, multiclass classification, continuous regression, and survival analysis. However, the audit diagnostics and simulation study presented here focus on binary classification. The permutation-gap test and delta LSI framework extend naturally to regression metrics such as RMSE and R^2, but these have not been evaluated empirically in the current manuscript.

Additional limitations are more specific to the current diagnostics. The univariate target scan is conservative by design and can miss indirect or multivariate proxies, while the stored supplementary results show that the multivariate scan may respond to general predictive signal rather than to leakage alone. Duplicate detection depends on the supplied feature space and the chosen similarity threshold. Association tables can reflect the intended split policy rather than a defect, as the case study illustrates for study holdout. Nested tuning does not yet support survival outcomes. Target scanning is unavailable for survival tasks. Exchangeability options `by_group` and `within_batch` are recorded by `delta_lsi()` but currently use the independent sign-flip engine rather than a specialized randomization scheme. The package supports multiclass modeling, but leakage diagnostics are more developed for binary classification than for other tasks.

9.3. Future extensions

These restrictions are software design choices rather than conceptual impossibilities, and they define concrete directions for further development. Future work could expand the leakage taxonomy to include external-knowledge leakage, improve multivariate target-leakage calibration, extend nested tuning to survival models, and add more formal summaries for duplicate-risk severity. Another useful extension would be tighter integration with manuscript generation, so that audit reports, scalar summaries, and figure artifacts can be linked automatically in literate-analysis pipelines.

10. Summary

`bioLeak` provides an integrated R workflow for preventing and diagnosing data leakage in biomedical machine learning. The package pairs leakage-aware resampling and guarded preprocessing with post hoc permutation tests, batch-association diagnostics, target-leakage scans, and duplicate detection. The `delta_lsi()` framework offers formal inference for comparing naive and guarded pipelines. The simulation study shows that the permutation-gap

test reliably detects several leakage mechanisms at moderate sample sizes, while the case study demonstrates how audit components interact on real multi-study transcriptomic data. Together, these tools make leakage assessment a documented, reproducible part of the modeling record.

11. Availability

As of March 13, 2026, `bioLeak` is available on CRAN as version 0.3.0. The development source repository is hosted at https://github.com/selcukorkmaz/bioLeak. Package documentation includes a README, a long-form vignette (`vignettes/bioLeak-intro.Rmd`), and function-level help pages installed with the package. The source repository also contains the paper materials used in this manuscript, including `paper/run_simulation.R`, `paper/run_supplementary.R`, `paper/run_casestudy.R`, `paper/run_delta_lsi.R`, `paper/compile_results.R`, `paper/generate_figures` and the stored result objects discussed in Sections 6 and 7. The package is released under the MIT license.

Computational details

The results in this paper were obtained using R 4.4.2 (**?**). R itself and all packages used are available from the Comprehensive R Archive Network (CRAN) at https://CRAN.R-project.org/.

The principal package is **bioLeak** 0.3.0. Key dependencies include **parsnip** 1.2.1 for model specification, **recipes** 1.1.0 for preprocessing pipelines, **glmnet** 4.1-8 for regularized regression, **pROC** 1.18.5 for AUC computation, and **future** 1.34.0 and **future.apply** 1.11.3 for parallel execution. Version numbers should be verified against the analysis environment using `sessionInfo()`.

Acknowledgments

The author thanks the anonymous reviewers for constructive feedback.

References

Ambroise C, McLachlan GJ (2002). "Selection Bias in Gene Extraction on the Basis of Microarray Gene-Expression Data." *Proceedings of the National Academy of Sciences*, **99**(10), 6562–6566. doi:10.1073/pnas.102102699.

Arlot S, Celisse A (2010). "A Survey of Cross-Validation Procedures for Model Selection." *Statistics Surveys*, **4**, 40–79. doi:10.1214/09-SS054.

Bates S, Hastie T, Tibshirani R (2024). "Cross-Validation: What Does It Estimate and How Well Does It Do It?" *Journal of the American Statistical Association*, **119**(546), 1434–1445. doi:10.1080/01621459.2023.2197686.

Beam AL, Kohane IS (2018). "Big Data and Machine Learning in Health Care." *JAMA*, **319**(13), 1317–1318. doi:10.1001/jama.2017.18391.

Bengio Y, Grandvalet Y (2004). "No Unbiased Estimator of the Variance of K-Fold Cross-Validation." *Journal of Machine Learning Research*, **5**, 1089–1105.

Bischl B, Sonabend R, Kotthoff L, Lang M (2024). *Applied Machine Learning Using mlr3 in R*. CRC Press.

Cawley GC, Talbot NLC (2010). "On Over-Fitting in Model Selection and Subsequent Selection Bias in Performance Evaluation." *Journal of Machine Learning Research*, **11**, 2079–2107.

Christen P (2012). *Data Matching: Concepts and Techniques for Record Linkage, Entity Resolution, and Duplicate Detection*. Springer. doi:10.1007/978-3-642-31164-2.

Davis SE, Greevy RA, Fonnesbeck C, Lasko TA, Walsh CG, Matheny ME (2024). "A Nonparametric Updating Method to Correct Clinical Prediction Model Drift." *Journal of the American Medical Informatics Association*, **31**(6), 1275–1282. doi:10.1093/jamia/ocae030.

Efron B (1979). "Bootstrap Methods: Another Look at the Jackknife." *The Annals of Statistics*, **7**(1), 1–26. doi:10.1214/aos/1176344552.

Efron B, Tibshirani RJ (1993). *An Introduction to the Bootstrap*. Chapman and Hall/CRC.

Gentleman R, Lang DT (2007). "Statistical Analyses and Reproducible Research." *Journal of Computational and Graphical Statistics*, **16**(1), 1–23. doi:10.1198/106186007X178663.

Golub TR, Slonim DK, Tamayo P, Huard C, Gaasenbeek M, Mesirov JP, Coller H, Loh ML, Downing JR, Caligiuri MA, Bloomfield CD, Lander ES (1999). "Molecular Classification of Cancer: Class Discovery and Class Prediction by Gene Expression Monitoring." *Science*, **286**(5439), 531–537. doi:10.1126/science.286.5439.531.

Hamdan S, Love BC, von dem Hagen EAH (2023). "Confound-Leaking: Regression Leak in Machine Learning Applied to Neuroimaging." 2306.12952.

Hastie T, Tibshirani R, Friedman J (2009). *The Elements of Statistical Learning: Data Mining, Inference, and Prediction*. 2nd edition. Springer. doi:10.1007/978-0-387-84858-7.

Huber PJ (1964). "Robust Estimation of a Location Parameter." *The Annals of Mathematical Statistics*, **35**(1), 73–101. doi:10.1214/aoms/1177703732.

Huber PJ, Ronchetti EM (2009). *Robust Statistics*. 2nd edition. Wiley. doi:10.1002/9780470434697.

James G, Witten D, Hastie T, Tibshirani R (2013). *An Introduction to Statistical Learning*. Springer. doi:10.1007/978-1-4614-7138-7.

Kapoor S, Narayanan A (2023). "Leakage and the Reproducibility Crisis in Machine-Learning-Based Science." *Patterns*, **4**(9), 100804. doi:10.1016/j.patter.2023.100804.

Kaufman S, Rosset S, Perlich C, Stitelman O (2012). "Leakage in Data Mining: Formulation, Detection, and Avoidance." *ACM Transactions on Knowledge Discovery from Data*, **6**(4), 1–21. doi:10.1145/2382577.2382579.

Kohavi R (1995). "A Study of Cross-Validation and Bootstrap for Accuracy Estimation and Model Selection." In *Proceedings of the 14th International Joint Conference on Artificial Intelligence (IJCAI)*, volume 14, pp. 1137–1145.

Kuhn M (2008). "Building Predictive Models in R Using the **caret** Package." *Journal of Statistical Software*, **28**(5), 1–26. doi:10.18637/jss.v028.i05.

Kuhn M, Wickham H (2020). "Tidymodels: A Collection of Packages for Modeling and Machine Learning Using Tidyverse Principles." URL https://www.tidymodels.org.

Lang M, Binder M, Richter J, Schratz P, Pfisterer F, Coors S, Au Q, Casalicchio G, Kotthoff L, Bischl B (2019). "**mlr3**: A Modern Object-Oriented Machine Learning Framework in R." *Journal of Open Source Software*, **4**(44), 1903. doi:10.21105/joss.01903.

Leskovec J, Rajaraman A, Ullman JD (2014). *Mining of Massive Datasets*. 2nd edition. Cambridge University Press.

Moscovich A, Rosset S (2022). "On the Cross-Validation Bias Due to Unsupervised Pre-Processing." *Journal of the Royal Statistical Society: Series B (Statistical Methodology)*, **84**(4), 1474–1502. doi:10.1111/rssb.12537.

Peng RD (2011). "Reproducible Research in Computational Science." *Science*, **334**(6060), 1226–1227. doi:10.1126/science.1213847.

R Core Team (2024). *R: A Language and Environment for Statistical Computing*. R Foundation for Statistical Computing, Vienna, Austria. URL https://www.R-project.org/.

Roberts DR, Bahn V, Ciuti S, Boyce MS, Elith J, Guillera-Arroita G, Hauenstein S, Lahoz-Monfort JJ, Schröder B, Thuiller W, Warton DI, Wintle BA, Hartig F, Dormann CF (2017). "Cross-Validation Strategies for Data with Temporal, Spatial, Hierarchical, or Phylogenetic Structure." *Ecography*, **40**(8), 913–929. doi:10.1111/ecog.02881.

Rosenblatt JD, Ben Simhon Y, Gal S, Gilron R (2024). "Practicalities of Resampling Methods for Brain Imaging: What Every Researcher Should Know." *NeuroImage*, **287**, 120520. doi:10.1016/j.neuroimage.2024.120520.

Sandve GK, Nekrutenko A, Taylor J, Hovig E (2013). "Ten Simple Rules for Reproducible Computational Research." *PLOS Computational Biology*, **9**(10), e1003285. doi:10.1371/journal.pcbi.1003285.

Sendak M, Elish MC, Gao M, Futoma J, Ratliff W, Nichols M, Bedoya A, Balu S, O'Brien C (2020). "The Human Body Is a Black Box: Supporting Clinical Decision-Making with Deep Learning." In *Proceedings of the Conference on Fairness, Accountability, and Transparency (FAT* '20)*, pp. 99–109. doi:10.1145/3351095.3372827.

Simon R, Radmacher MD, Dobbin K, McShane LM (2003). "Pitfalls in the Use of DNA Microarray Data for Diagnostic and Prognostic Classification." *Journal of the National Cancer Institute*, **95**(1), 14–18. doi:10.1093/jnci/95.1.14.

Soneson C, Delorenzi M (2014). "Batch Effect Confounding Leads to Strong Bias in Performance Estimates Obtained by Cross-Validation." *PLOS ONE*, **9**(6), e100335. doi:10.1371/journal.pone.0100335.

Stone M (1974). "Cross-Validatory Choice and Assessment of Statistical Predictions." *Journal of the Royal Statistical Society: Series B (Methodological)*, **36**(2), 111–147. doi:10.1111/j.2517-6161.1974.tb00994.x.

Subramanian J, Simon R (2013). "Overfitting in Prediction Models – Is It a Problem Only in High Dimensions?" *Contemporary Clinical Trials*, **36**(2), 636–641. doi:10.1016/j.cct.2013.06.011.

Tarca AL, Carey VJ, Chen Xw, Romero R, Drăghici S (2007). "Machine Learning and Its Applications to Biology." *PLOS Computational Biology*, **3**(6), e116. doi:10.1371/journal.pcbi.0030116.

Valavi R, Elith J, Lahoz-Monfort JJ, Guillera-Arroita G (2019). "**blockCV**: An R Package for Generating Spatially or Environmentally Separated Folds for k-Fold Cross-Validation of Species Distribution Models." *Methods in Ecology and Evolution*, **10**(2), 225–232. doi:10.1111/2041-210X.13107.

van de Mortel LA, van Wingen GA (2025). "Data Leakage Inflates Prediction Performance in Connectome-Based Machine Learning Models." *Nature Communications*, **16**, 1106. doi:10.1038/s41467-025-56082-4.

Varma S, Simon R (2006). "Bias in Error Estimation When Using Cross-Validation for Model Selection." *BMC Bioinformatics*, **7**, 91. doi:10.1186/1471-2105-7-91.

Wenger SJ, Olden JD (2012). "Assessing Transferability of Ecological Models: An Underappreciated Aspect of Statistical Validation." *Methods in Ecology and Evolution*, **3**(2), 260–267. doi:10.1111/j.2041-210X.2011.00170.x.

Winkler WE (1999). "The State of Record Linkage and Current Research Problems." *Statistical Research Report Series RR99/04*, U.S. Bureau of the Census.

A. Delta LSI Case Study Output

The following is the full `summary()` output of the `LeakDeltaLSI` object returned by `delta_lsi()` when comparing the naive and guarded case-study pipelines described in Section 7. Both fits share the same batch-blocked fold structure with 20 repeats, enabling paired repeat-level inference.

```
Delta LSI Summary
=================
Metric:            auc
Exchangeability:   iid
Inference tier:    A_full_inference
R_eff:             20

Point estimates:
  delta_metric (mean):  0.085
  delta_lsi (Huber):    0.085

Sign-flip test:
  p-value:              < 0.001
```

The inference tier is `A_full_inference` because 20 repeats are sufficient for the sign-flip randomization test and BCa bootstrap intervals (≥ 20 repeats yield tier A full inference).

B. Delta LSI Simulation

The script `paper/run_delta_lsi.R` includes a small simulation to evaluate the operating characteristics of `delta_lsi()`. The design uses $n = 200$, $p = 20$, 20 repeats of five-fold CV, and 20 seeds per condition. Two conditions are evaluated:

- **Null condition**: both pipelines process the same outcome, but each receives a different set of three pure-noise features (uncorrelated with y). Both arms share the same fold structure, enabling paired inference at tier A. Because the noise features are drawn independently, the per-repeat deltas fluctuate around zero rather than equaling zero identically—this exercises the sign-flip machinery non-trivially. Under this design the observed rejection rate was 15.0% at $\alpha = 0.05$ (mean $\widehat{\Delta}_{\text{metric}} = -0.0007$).

- **Alternative condition**: the naive pipeline includes a `peek_norm` leakage feature. The rejection rate was 100.0% with a mean $\widehat{\Delta}_{\text{metric}} = 0.1927$.

Results are stored in `paper/sim_results/delta_lsi_sim.rds`.

Affiliation:

Selçuk Korkmaz
Department of Biostatistics and Medical Informatics
Faculty of Medicine
Trakya University
Balkan Yerleşkesi
22030 Edirne, Türkiye
E-mail: selcukorkmaz@gmail.com
URL: https://selcukorkmaz.github.io/

Journal of Statistical Software http://www.jstatsoft.org/
published by the Foundation for Open Access Statistics http://www.foastat.org/
MMMMMM YYYY, Volume VV, Issue II *Submitted:* yyyy-mm-dd
doi:10.18637/jss.v000.i00 *Accepted:* yyyy-mm-dd